

Operating Systems – I (CS3510)

Assignment – II (Loading and Unloading Kernel modules)

Aim – To write a Linux kernel module that displays a list of all current running tasks in a Linux system. (Alternative command - “ps -el” on Linux Bash terminal).

Kernel Module : They are pieces of code that can be loaded and unloaded into the kernel upon demand. They extend the functionality of the kernel without rebooting the system.

Program Skeleton :

#include “header files”

module_init() function definition {

*for_each_process(p)
 print task_details;*

}

module_exit() function definition {

print exit message;

}

module_init();

module_exit();

module_license()

module_author()

module_description()

Header files included :

```
#include <linux/module.h>           // Needed to write kernel module
#include <linux/kernel.h>           // Needed for KERN_INFO
#include <linux/sched.h>            // Needed for task_struct
#include <linux/init.h>             // Needed for module_init() and module_exit() macros
#include <linux/dcache.h>
```

“linux/module.h” is needed to write the code into module entry point.

“kernel.h” is required for writing optional prefix strings called **Loglevel** inside the **printk** function like **KERN_INFO**, **KERN_CRIT**, **KERN_ALERT**, for informational, critical conditions report and prompt for action must be taken immediately purposes.

“sched.h” has the macro defined for “**for_each_process**”, as well as the **struct – task_struct**, defined in it.

“**init.h**” has the definitions of the macros **module_init(f)** and **module_exit(fe)** as given below.

INIT and CLEANUP functions

We can rename the init and cleanup functions of our kernel module. This is done with the **module_init()** and **module_exit()** macros. These macros are defined in **linux/init.h**.

We define our init and cleanup functions before calling the macros, as **start_module()** and **end_module()**, in order to avoid any compilation errors.

Loading the Kernel module into the Linux kernel

When **module_init()** is invoked (upon running **sudo insmod ./*.ko**), the code written in the function **start_module()** gets executed, where using the macro “**for_each_process(struct task_struct *p)**”, we print each task’s name (executable name), its PID, and the state associated with it.

for_each_process() macro allows iteration over all current tasks in the system. We use **printk** function to print these contents in the kernel log buffer. We can specify optional **LogLevel** strings to provide special status to the content being printed. The default if not specified is **KERN_WARNING**. **KERN_INFO** is the one we are providing for printing welcome and exit messages at the time of module loading and unloading respectively.

Unloading the Kernel module from the Linux kernel

When the user wants to remove this loaded module into the kernel by running the command **sudo rmmod *modulename*** (which *unloads* the kernel module from the Linux kernel), the function called in the macro **module_exit()** will be executed, at module removal time. Therefore, to print an exit message, we define a function “**end_module()**”, which is invoked at the time of unloading the kernel module/invoking of the macro **module_exit()**.

The argument to both **start_module()** and **end_module()** functions is *void* here.

Once the code is executed, and the kernel module is loaded, the output of **printk** is inserted in the **kernel log buffer**. This output can be viewed by entering the **dmesg** command on bash terminal which shows the output of **printk**, and the name of the module inserted can be viewed by entering the **lsmod** command, which generally appears on the top of the list of the other kernel modules.

lsmod Output : Here, the name of the module is **task_list** -

```
saurabh_cs_iith@saurabhcsiith-Inspiron-3542:~/Desktop/Assgn2-CS14BTECH11031/src$ lsmod
Module                  Size  Used by
Assgn2Src_CS14BTECH11031 16384  0
```

dmesg Output : This is a partial snapshot of some of the tasks, with their *name*, *state* and *PID* printed as shown below -

```

[20279.931071] Task-Name : systemd -|- Task-State : 1 -|- PID : 1
[20279.931074] Task-Name : kthreadd -|- Task-State : 1 -|- PID : 2
[20279.931077] Task-Name : ksoftirqd/0 -|- Task-State : 1 -|- PID : 3
[20279.931079] Task-Name : kworker/0:0H -|- Task-State : 1 -|- PID : 5
[20279.931081] Task-Name : rcu_sched -|- Task-State : 1 -|- PID : 7
[20279.931083] Task-Name : rcu_bh -|- Task-State : 1 -|- PID : 8
[20279.931085] Task-Name : migration/0 -|- Task-State : 1 -|- PID : 9
[20279.931087] Task-Name : watchdog/0 -|- Task-State : 1 -|- PID : 10
[20279.931090] Task-Name : watchdog/1 -|- Task-State : 1 -|- PID : 11
[20279.931092] Task-Name : migration/1 -|- Task-State : 1 -|- PID : 12
[20279.931094] Task-Name : ksoftirqd/1 -|- Task-State : 1 -|- PID : 13
[20279.931096] Task-Name : kworker/1:0H -|- Task-State : 1 -|- PID : 15
[20279.931099] Task-Name : watchdog/2 -|- Task-State : 1 -|- PID : 16
[20279.931101] Task-Name : migration/2 -|- Task-State : 1 -|- PID : 17
[20279.931103] Task-Name : ksoftirqd/2 -|- Task-State : 1 -|- PID : 18
[20279.931105] Task-Name : kworker/2:0H -|- Task-State : 1 -|- PID : 20
[20279.931107] Task-Name : watchdog/3 -|- Task-State : 1 -|- PID : 21
[20279.931109] Task-Name : migration/3 -|- Task-State : 1 -|- PID : 22
[20279.931111] Task-Name : ksoftirqd/3 -|- Task-State : 1 -|- PID : 23
[20279.931114] Task-Name : kworker/3:0H -|- Task-State : 1 -|- PID : 25
[20279.931116] Task-Name : kdevtmpfs -|- Task-State : 1 -|- PID : 26
[20279.931119] Task-Name : netns -|- Task-State : 1 -|- PID : 27
[20279.931121] Task-Name : perf -|- Task-State : 1 -|- PID : 28
[20279.931123] Task-Name : khungtaskd -|- Task-State : 1 -|- PID : 29
[20279.931125] Task-Name : writeback -|- Task-State : 1 -|- PID : 30

```

sudo insmod *module_name* **Output** : Message “Kernel module successfully loaded !” is printed once the kernel module is loaded.

```

[20279.931068] Kernel module successfully loaded !
[20279.931071] Task-Name : systemd -|- Task-State : 1 -|- PID : 1
[20279.931074] Task-Name : kthreadd -|- Task-State : 1 -|- PID : 2

```

sudo rmmod *module_name* **Output** : Message “Kernel module successfully unloaded !” is printed in the log buffer, once the module is successfully unloaded.

```

[20429.849228] Kernel module successfully unloaded !
saurabh_cs_iith@saurabhcsiith-Inspiron-3542:~/Desktop/Assgn2-CS14BTECH11031/src$

```

Note that **insmod** and **rmmod** commands require root privileges. The output of the **dmesg** command can be confirmed by running the **ps -el** command.

The user can compile and run the program by going into the “src” directory and running the **make** command. Then **insmod** to load the kernel module, **lsmod** to check the name of module, **dmesg** to scan the output given by *printk()*, and finally **rmmod** to remove the module from the Linux kernel.

Details regarding execution and loading/unloading is described in the readme file in detail.